

Probabilistic programming languages for Bayesian inference in R

Måns Magnusson

StiMa, Linköping University

- BUGS (**B**ayesian inference **U**sing **G**ibbs **S**ampling)
- JAGS (**J**ust **A**nother **G**ibbs **S**ampler)
- Stan (**S**tanislav Ulam)

When not to use

- Big Data
- Big Models

Why JAGS?

- Easy to use (BUGS type)
- Robust (version 4.x.x)
- Wrappers for Python and R
- No C++ compiler needed
- Good starting point

JAGS website

Why Stan?

- Easy to use (BUGS type)
- Effective samplers/fast
- Integrates nice with RStudio
- Wrappers from Python, R, Matlab, Stata, Julia
- Automatic/Black box variational inference
- (Really) good documentation

mc-stan.org

- Bernoulli-Beta (Coin flips, conjugate)

- Model:

$$y_i \sim \text{Bern}(p)$$

- Prior:

$$p \sim \text{Beta}(a, b)$$

- Poisson regression (Number of roaches caught in buildings, non-conjugate)

- Model:

$$y_i \sim \text{Poisson}(\log(\text{exposure}) + \beta_1 + \beta_2 \cdot \text{treatment} + \beta_3 \cdot \text{senior})$$

- Prior:

$$\beta \sim \text{Norm}(0, 1000)$$

- Needs C++ compiler - good installation instructions [here](#)
- Six parts in a Stan model:
 - data
 - transformed data
 - parameters
 - transformed parameters
 - model*
 - generated quantities

- Read in data (for example from R or Python) once
 - Only variable declarations
 - A lot of different data types
 - `int`, `real`, `vector`, `arrays`, `matrix` and more Stan specific (?) data types as `cholesky_factor_cov` and `unit_vector`

Example of the data block

```
data {  
  int<lower=0> N; # The number of observations  
  int<lower=0> y;  
  vector[N] exposure2;  
  vector[N] senior;  
  vector[N] treatment;  
}
```


- Variable declarations **and** statements (done once)
 - See chapter V in the documentation for all functions that can be used.

Example of the transformed data block

```
transformed data {  
  vector[N] log_expo;  
  log_expo <- log(exposure2);  
}
```

- Parameters (that should be sampled)
 - Parameter declarations only.

Example of the `parameters` block

```
parameters {  
  vector[3] beta;  
}
```

- Parameter declarations **and** statements
 - The transformations is done in each sampling step

Example of the transformed parameters block (not example model)

```
transformed parameters {  
  real<lower=0> sigma;  
  sigma <- 1.0 / sqrt(tau);  
}
```

- Declare the priors and data with sampling statements ~
 - Distributions can be found in chapter VI and VII in the documentation

Example of the model block

```
model {  
  // Priors  
  beta ~ normal(0.0, 1000.0);  
  # Model  
  y ~ poisson_log(log_expo + beta[1] + beta[2] * treatment +  
beta[3]*senior);  
}
```

- Computations after the sampling has been done, used for
- Is used for:
 - model checking
 - predictive distributions for new data
 - applying full Bayesian decision theory
 - transforming parameters for reporting, etc

Example of the generated quantities block (not example model)

```
generated quantities {  
  real my_weight_pred;  
  my_weight_pred <- alpha + beta * MySHeight +  
  normal_rng(0,sigma);  
}
```

How to specify a model in JAGS

- One parts in a JAGS model:

- model*

Example of the model block

```
model {  
  # Model  
  for( i in 1:N){  
    y[i] ~ dpois(lambda[i])  
    log(lambda[i]) <- log(exposure2[i]) + beta1 + beta2 *  
treatment[i] + beta3 * senior[i])  
  }  
  # Priors  
  beta1 ~ dnorm(0.0, 0.0001)  
  beta2 ~ dnorm(0.0, 0.0001)  
  beta3 ~ dnorm(0.0, 0.0001)  
}
```

Demonstration